BiobankCloud: Secure Storage, Sharing, and Processing of NGS Data

Alysson Bessani[5], Marc Bux[2], Joergen Brandt[2], Jim Dowling[1], Ali Gholami[1], Micheal Hummel[4], Mahmoud Ismail[1], Ulf Leser[2], Jan-Eric Litton[3], Roxanna Martinez[3], Salman Niazi[1]

[1] KTH - Royal Institute of Technology,
{jdowling, gholami, maism, smkniazi}@kth.se
[2] Humboldt University
{leser, bux, joergen.brandt}@informatik.hu-berlin.de
[3] Karolinska Institute
{Jan-Eric.Litton, Roxanna.Martinez}@ki.se
[4] Charite
{Michael.Hummel}@charite.de
[5] University of Lisbon
{bessani}@di.fc.ul.pt

Abstract. Biobanks store and catalogue human biological material that is increasingly becoming digitized using technologies such as whole genome sequencing. There is, however, a growing Biobank bottleneck: a lack of open-source software systems that can securely scale to handle the coming wave of genomic data from Next-Generation Sequencing machines. In the BiobankCloud project, we are building an open-source Hadoop-based platform for the secure storage, sharing, and processing of genomic data. We have extended Hadoop to include support for multi-tenancy (studies), reduced data sizes with erasure coding, support for extensible and consistent metadata, a scalable scientific workflow framework, and support for securely sharing data across Hadoop instances. We demonstrate the capabilities of BiobankCloud in a Laboratory Information Managment System with support for 2-factor authentication and UI-support to all major services.

1 Introduction

Biobanks store and catalog human biological material from identifiable individuals for both clinical and research purposes. Recent advances in Next-Generation Sequencing (NGS) technology has meant that there is an increasing demand to sequence the human biological material stored in Biobanks. Both research projects and clinical health-care systems use Biobanks to store samples that are sequenced using techniques from genotyping to whole-genome sequencing (WGS). Biobanks' computer systems have traditionally managed only metadata associated with samples, such as pseudo-identifiers for patients, sample collection information, study infromation, and data concerning samples. Alongside

this metadata, we now increasingly need to store genomic data, which requires anything from MBs (genotyping) up to several hundred GBs (WGS) of data per sample. The data storage requirements for large-scale WGS sequencing projects, such as the 100,000 genomes project by Genomics England, now scale to many PBs, and are so large that researchers are looking at cost-effective Big Data solutions based on commodity hardware, such as Hadoop. When data volumes grow past several TBS, traditional database technology (including sharded relational databases) is no longer viable, as more data needs to be moved from storage to compute nodes that is possible with current networking technology. The solution provided by platforms such as Hadoop is to move computation to where the data is located, exploiting the principle of *data locality*. That is, jobs are parallelized across many nodes and each job loads its (large) input data primarily from disk subsystems, while network I/O is used transfer in subsequent processing and sorting steps on, typically, smaller data volumes relative to the input data size.

In this paper, we introduce BiobankCloud, an integrated platform, based on Hadoop, for the secure storage, processing, and sharing of genomic data and associated metadata. BiobankCloud is a platform-as-a-service that can be automatically deployed on public clouds, private clouds or bare-metal servers. As part of BiobankCloud, we also provide a Laboratory Information Management Service (LIMS) as software-as-a-service (SaaS). The LIMS has an integrated User Interface (UI) for authenicating/authorizing users, managing data, designing and searching for metadata, and support for running workflows and analysis jobs on Hadoop. The LIMS hides much of the complexity of the Hadoop backend, and supports multi-tenancy through first-class support for *Studies*, *SampleCollections* (DataSets), *Samples*, and *Users*.

In BiobankCloud, we have developed our own Hadoop distribution, Hadoop Open Platform-as-a-Service (Hops), with improved scalability and customizability properties, enabled by a new metadata storage system based on a distributed in-memory database.

Existing scientific workflow management systems are typically custom-built and have not been designed for data parallel processing with data locality. For performance reasons, their architectures and file formats are typically flat (monolithic). None of the main file formats in genomics (fastq, BAM/SAM, CRAM, and VCF) have support for data parallel processing because of their assumption of centralized metadata (e.g., in the file header). In BiobankCloud, we provide a scientific workflow management system, SaasFee, as a native YARN/Hadoop 2nd-level scheduler that provides a bridge between support for existing file formats and data parallel processing. SaasFee can both speedup workflows, such as a ??x speedup for NGS varient calling on XX machines, and scale-out to support larger clusters.

2 Related work

There are a couple of frameworks for data parallel processing of genomic data. Adam, Halvade, Seal, PigSeq, Spork?

There's the big Alvados? project in Harvard.

In security, Hadoop has support for Apache Ranger (attribute-based access control), Apache Sentry (Databases, RBAC), and Apache ?? (RBAC for REST APIs).

Nothing happening in Biobanking?
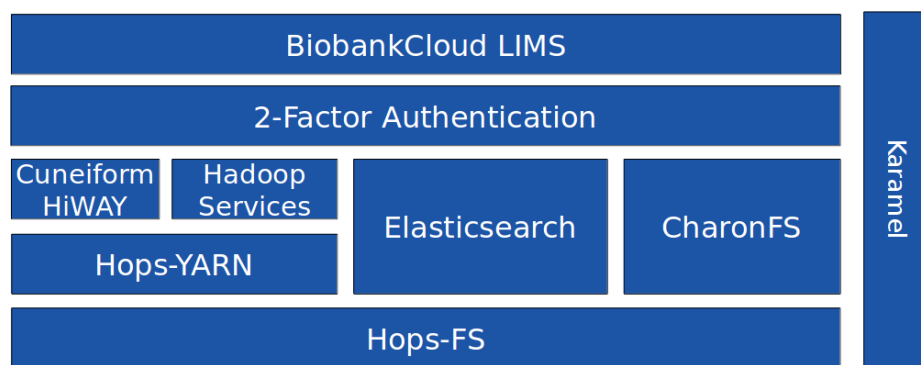
3 BiobankCLoud Regulatory Framework

4 LIMS

Fig. 1: LIMS

In BiobankCloud, we are using our distribution of the Hadoop Filesystem (HDFS), HopsFS, to store large genomic files. HopsFS scales to store 100s of millions of files. There is a need to organize the files in such a manner that they can easily be logically grouped and searched. Typically, such functionality requires metadata for the files and directories, such as filename and last accessed time. For biological samples, much more extensive metadata is required for files that relate to biological samples. We need information such as the sample collection it belongs to, the type of sample, and donor information. BiobankCloud enables Biobankers who are not programmers to design such metadata, link it to samples, and then edit sample metadata using UI support. The metadata is transparently indexed to enable free-text searching for samples, sample collections, and studies. Our solution is based on the industry-standard Elasticsearch platform. Importantly, we guarantee the metadata's integrity by implementing an eventually consistent replication model that asynchronously copies updates to metadata from our distributed database to Elasticsearch. The distributed database uses foreign keys to guarantee the integrity of the metadata and the referent genomic data. Mutations to the metadata or removal of the sample will mutate or remove the metadata in Elasticsearch within seconds.

Biobankers can use a web application to design new tables that are transparently added to the database and have foreign-key constraints on existing files or directories in the system, thus maintaining the integrity of the metadata. Our system automatically exports metadata to Elasticsearch from where it can be searched using freetext by the user. We support both the automated indexing of files and directories in Elasticsearch as well as custom-designed metadata.

BiobankCloud introduces **DataSets** as a new abstraction to Hadoop, where a DataSet consists of a related group of directories, files, and extended metadata. DataSets can be indexed and searched and are the basic unit of data management in BiobankCloud; all user-generated files or directories belong to a single DataSet. In Biobanking, a sample collection would be a typical example of a DataSet.

To allow for access control of users to DataSets, which is not inherent in the DataSet concept, we introduce the notion of **Studies**. A Study is a grouping of researchers and DataSets with role-based access control where different researchers can be given different access rights to DataSets. The basic user roles we provide reflect the European Data Protection Directive, with a DataOwner (data controller) and a DataScientist (data processer). DataSets can be shared between Studies (when the necessary security, legal, and ethical conditions for sharing are in place). In BiobankCloud, we use the access control mechanism of HopsFS to implement the Study- and DataSet-based authorization model.

5 Security Model

The BioBankCloud environment deploys strong security features for concerns such as Confidentiality, Integrity and Non-repudiation [?] of data access. This includes authentication, authorization , and auditing. The system allows defining different roles with different access privileges. In designing the system, we applied the Cloud Privacy Threat Modeling [?] approach to identify the privacy requirements of processing sensitive biomedical data.

Figure **??** shows the different components of the employed security mechanisms. All BioBankCloud services are protected behind the firewall and only accessible through the secure interfaces over HTTPS Channels.

5.1 2-Factor Authentication

The authentication services map the person accessing the platform to a user identity. We provide two-factor authentication using smart mobile devices or Yubikey hardware tokens to support different groups of users.

5.2 Access Control

The accessc control component ensures authorized access to genomics data as internal objects or different services within the platform. This is accomplished through a role-based access control (RBAC) model and a data access approach
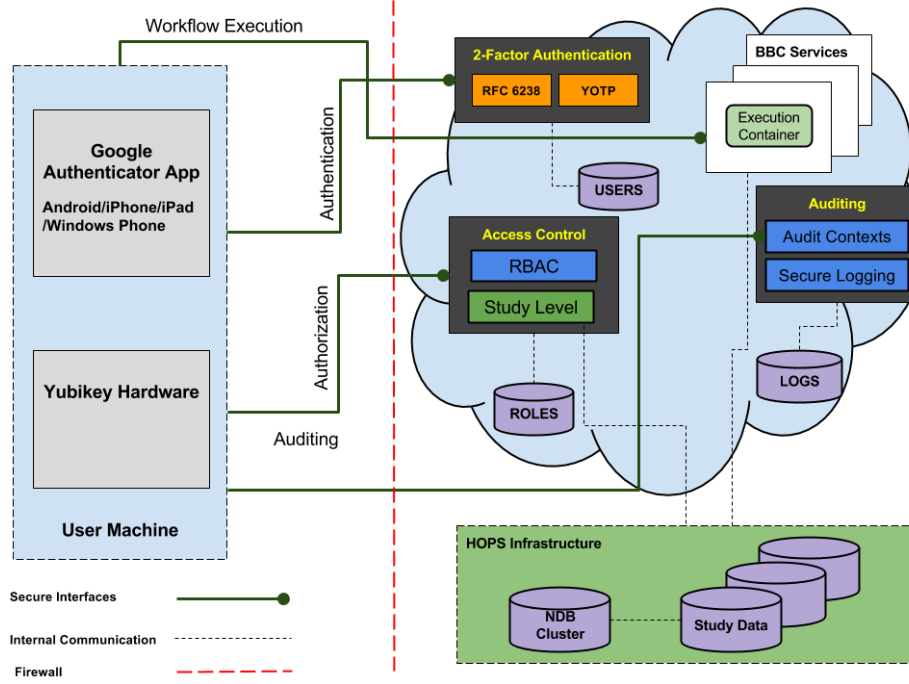
Fig. 2: Security Architecture of the BiobankCloud

using the ownership of data on the HDFS, for example, a data owner adds/revokes members to a study and assigns privileges to access the study.

5.3 HDFS Supported Authorization

5.4 Auditing Service

Finally, the auditing service enables the platform administrator or an external auditor to discover the history of accessing the platform to detect any violation to a policy. It includes several contexts such as role, account, study, and login audits. The secure login service assures that actions that are taken by the users are registered for tracing and auditing purposes. Each log event contains informat ion such as initiator, target, IP/MAC addresses, timestamp, action, and outcome.

6 Hadoop Open Platform-as-a-Service (Hops)

We have moved metadata from the heap of a Java Virtual Machine to an in-memory, no-shared-state, distributed database, called MySQL Cluster [?].

```
HDFS
  ├→user
  ├→tmp
  └→studies
        ├→study₁
        │     └→datasetₐ
        ├→study₂
        │     └→dataset_B
        ├→...
        └→study_N
              ├→dataset_C
              └→dataset_D
```

Fig. 3: Filesystem structure in HDFS containing Studies and DataSets.

We ensure the consistency of the filesystem metadata by implementing serialized transactions on well-ordered operations on metadata [**?**]. We do not need Zookeeper, as we have developed a leader-election service based on the database [**?**].

7 Sassfee

8 Sharing Data Between Clusters with CharonFS

9 Conclusions

In this paper, we introduced the BiobankCloud platform, a Hadoop-based platform that provides a number of features necessary for Biobanks to adopt Hadoop-based solutions for managing NGS data. Critical safety features that we introduced for managing sensitive data include multi-tenancy for the isolation of Study data and secure access through 2-factor authentication. Next-generation data management systems for NGS data must be massively scalable. We introduced our scalable storage service, HopsFS, our processing framework, Hops-YARN, and our framework for scalable bioinformatics workflows, Saasfee. We also showed the integration of explicity metadata design and management in the platform, ensuring the integrity of metadata and supporting free-text search using extended metadata. Finally, in CharonFS, we showed how we can leverage public clouds to share data securely between clusters.

10 Acknowledgements